

assignment_1

Oskar Våle

2024-09-06

```
knitr::opts_chunk$set(echo = TRUE)
library(reticulate)
use_python("/usr/bin/python3")
```

Power consumption dataset

The dataset I chose is the “Individual Household Electric Power Consumption” dataset from the UCI machine learning repository. The dataset contains measurements of electric power consumption in a single household over a period of almost 4 years. The data includes variables such as global active power, global reactive power, voltage, and global intensity, as well as sub-metering values corresponding to different household areas like the kitchen, laundry room, and other appliances. Additionally, the data includes timestamps that have been used to extract the time of day and the month in which the data was recorded.

This is the variable information from UCI’s page: 1.date: Date in format dd/mm/yyyy 2.time: time in format hh:mm:ss 3.global_active_power: household global minute-averaged active power (in kilowatt) 4.global_reactive_power: household global minute-averaged reactive power (in kilowatt) 5.voltage: minute-averaged voltage (in volt) 6.global_intensity: household global minute-averaged current intensity (in ampere) 7.sub_metering_1: energy sub-metering No. 1 (in watt-hour of active energy). It corresponds to the kitchen, containing mainly a dishwasher, an oven and a microwave (hot plates are not electric but gas powered). 8.sub_metering_2: energy sub-metering No. 2 (in watt-hour of active energy). It corresponds to the laundry room, containing a washing-machine, a tumble-drier, a refrigerator and a light. 9.sub_metering_3: energy sub-metering No. 3 (in watt-hour of active energy). It corresponds to an electric water-heater and an air-conditioner.

I decided to make one-hot encoded columns for whether a given datapoint is recorded in the morning, afternoon, evening or night instead of training on the ‘Time’ column. My hunch is that the model will perform better this way. I also made a new column ‘Month’, and let this one just be integers 1 - 12. Due to temperature, holidays and other factors, energy consumption is likely to vary from month to month. There is no point in training on days of the month though, as, except for weekends and holidays, there is no reason one day should have more consumption than the next.

```
import pandas as pd
from ucimlrepo import fetch_ucirepo
```

```
# Fetch dataset
individual_household_electric_power_consumption = fetch_ucirepo(id=235)
```

```
## /home/oskarva/.local/lib/python3.10/site-packages/ucimlrepo/fetch.py:97: DtypeWarning: Columns (2,3,
## df = pd.read_csv(data_url)
```

```
# Data (as pandas dataframes)
X = individual_household_electric_power_consumption.data.features
y = individual_household_electric_power_consumption.data.targets
```

```

# Variable information
print(individual_household_electric_power_consumption.variables)

##           name      role      type ... description units missing_values
## 0           Date  Feature      Date ...      None  None          no
## 1           Time  Feature  Categorical ...      None  None          no
## 2  Global_active_power  Feature  Continuous ...      None  None          no
## 3  Global_reactive_power  Feature  Continuous ...      None  None          no
## 4           Voltage  Feature  Continuous ...      None  None          no
## 5  Global_intensity  Feature  Continuous ...      None  None          no
## 6  Sub_metering_1  Feature  Continuous ...      None  None          no
## 7  Sub_metering_2  Feature  Continuous ...      None  None          no
## 8  Sub_metering_3  Feature  Continuous ...      None  None          no
##
## [9 rows x 7 columns]

# Combine features and targets for easier manipulation
data = pd.concat([X, y], axis=1)
#NOTE: The above code is taken from UCI's "import in python" function. This

# Convert 'Date' and 'Time' into a single datetime column
data['Datetime'] = pd.to_datetime(data['Date'] + ' ' + data['Time'], format='%d/%m/%Y %H:%M:%S')

# Function to categorize time of day
def categorize_time_of_day(hour):
    if 6 <= hour < 12:
        return 'Morning'
    elif 12 <= hour < 18:
        return 'Afternoon'
    elif 18 <= hour < 24:
        return 'Evening'
    else:
        return 'Night'

# Apply the function to create a new column 'Time_of_Day'
data['Time_of_Day'] = data['Datetime'].dt.hour.apply(categorize_time_of_day)

# Extract the month and create a new column 'Month'
data['Month'] = data['Datetime'].dt.month

# Drop the original Date, Time, and Datetime columns if not needed
data = data.drop(columns=['Date', 'Time', 'Datetime'])

# As the output shows, some of the data is still objects. We therefore need to
# convert it to numerical values.
print(data.dtypes)

## Global_active_power      object
## Global_reactive_power    object
## Voltage                  object
## Global_intensity         object
## Sub_metering_1           object
## Sub_metering_2           object
## Sub_metering_3          float64
## Time_of_Day              object

```

```

## Month                                int64
## dtype: object

cols_to_convert = ['Global_active_power', 'Global_reactive_power', 'Voltage',
                   'Global_intensity', 'Sub_metering_1', 'Sub_metering_2']

for col in cols_to_convert:
    data[col] = pd.to_numeric(data[col], errors='coerce')

#Check for NaN values
print(data.isna().sum())

```

```

## Global_active_power      25979
## Global_reactive_power    25979
## Voltage                  25979
## Global_intensity         25979
## Sub_metering_1          25979
## Sub_metering_2          25979
## Sub_metering_3          25979
## Time_of_Day              0
## Month                    0
## dtype: int64

# As the output shows, some of the data is still objects. We therefore need to
# convert it to numerical values.
print(data.dtypes)

```

```

## Global_active_power      float64
## Global_reactive_power    float64
## Voltage                  float64
## Global_intensity         float64
## Sub_metering_1          float64
## Sub_metering_2          float64
## Sub_metering_3          float64
## Time_of_Day              object
## Month                    int64
## dtype: object

```

```

# Drop rows that contain NaN values
data.dropna(axis=0, inplace=True)

# Perform one-hot encoding on the 'Time_of_Day' column, as mentionet above.. This
# is due to the fact that linear regression can only handle numerical values.
# I could also have had a single Time_of_Day column with morning=1,
# afternoon = 2, evening = 3, night = 4, but this implies that morning is far from
# night (4-1=3), which could negatively impact the model.
data = pd.get_dummies(data, columns=['Time_of_Day'], drop_first=False)

```

Including Plots

```

summary_stats = data.describe(include="all")

# Generating summary statistics for categorical variables (Time_of_Day after one-hot encoding)
categorical_summary = data[['Time_of_Day_Morning', 'Time_of_Day_Afternoon', 'Time_of_Day_Evening', 'Time_of_Day_Night']]

```

```
# Combine the summaries
```

```
pd.set_option('display.max_columns', None)
```

```
summary_stats_combined = summary_stats.append(categorical_summary)
```

```
## <string>:1: FutureWarning: The frame.append method is deprecated and will be removed from pandas in 1.0.0
```

```
summary_stats_combined.loc['count'] = len(data)
```

```
# Display the summary statistics
```

```
print(summary_stats_combined)
```

```
##      Global_active_power  Global_reactive_power      Voltage \
## count      2.049280e+06      2.049280e+06  2.049280e+06
## mean      1.091615e+00      1.237145e-01  2.408399e+02
## std      1.057294e+00      1.127220e-01  3.239987e+00
## min      7.600000e-02      0.000000e+00  2.232000e+02
## 25%      3.080000e-01      4.800000e-02  2.389900e+02
## 50%      6.020000e-01      1.000000e-01  2.410100e+02
## 75%      1.528000e+00      1.940000e-01  2.428900e+02
## max      1.112200e+01      1.390000e+00  2.541500e+02
## sum              NaN              NaN              NaN
##
##      Global_intensity  Sub_metering_1  Sub_metering_2  Sub_metering_3 \
## count      2.049280e+06      2.049280e+06      2.049280e+06      2.049280e+06
## mean      4.627759e+00      1.121923e+00      1.298520e+00      6.458447e+00
## std      4.444396e+00      6.153031e+00      5.822026e+00      8.437154e+00
## min      2.000000e-01      0.000000e+00      0.000000e+00      0.000000e+00
## 25%      1.400000e+00      0.000000e+00      0.000000e+00      0.000000e+00
## 50%      2.600000e+00      0.000000e+00      0.000000e+00      1.000000e+00
## 75%      6.400000e+00      0.000000e+00      1.000000e+00      1.700000e+01
## max      4.840000e+01      8.800000e+01      8.000000e+01      3.100000e+01
## sum              NaN              NaN              NaN              NaN
##
##      Month  Time_of_Day_Afternoon  Time_of_Day_Evening \
## count      2.049280e+06      2.049280e+06      2.049280e+06
## mean      6.454433e+00      2.498970e-01      2.504031e-01
## std      3.423209e+00      4.329533e-01      4.332453e-01
## min      1.000000e+00      0.000000e+00      0.000000e+00
## 25%      3.000000e+00      0.000000e+00      0.000000e+00
## 50%      6.000000e+00      0.000000e+00      0.000000e+00
## 75%      9.000000e+00      0.000000e+00      1.000000e+00
## max      1.200000e+01      1.000000e+00      1.000000e+00
## sum              NaN      5.121090e+05      5.131460e+05
##
##      Time_of_Day_Morning  Time_of_Day_Night
## count      2.049280e+06      2.049280e+06
## mean      2.496731e-01      2.500268e-01
## std      4.328239e-01      4.330283e-01
## min      0.000000e+00      0.000000e+00
## 25%      0.000000e+00      0.000000e+00
## 50%      0.000000e+00      0.000000e+00
## 75%      0.000000e+00      1.000000e+00
## max      1.000000e+00      1.000000e+00
## sum      5.116500e+05      5.123750e+05
```